

Coronavirus World Data Analysis

KATE expects your code to define variables with specific names that correspond to certain things we are interested in.

KATE will run your notebook from top to bottom and check the latest value of those variables, so make sure you don't overwrite them.

- Remember to uncomment the line assigning the variable to your answer and don't change the variable or function names.
- Use copies of the original or previous DataFrames to make sure you do not overwrite them by mistake.

You will find instructions below about how to define each variable.

Once you're happy with your code, upload your notebook to KATE to check your feedback.

First of all, run the following cell to:

- import pandas with an alias of pd
- read a CSV containing the data to work with
- convert the date column to the datetime format
- create a DataFrame df containing the data for only **1st July 2020**
- take a look at the first few rows of the DataFrame

In [1]: `import pandas as pd`

```
data = pd.read_csv('data/owid-covid-data.csv')
data['date'] = pd.to_datetime(data['date'])
df = data[data['date'] == '2020-07-01']

df.head()
```

Out[1]:

	iso_code	continent	location	date	total_cases	new_cases	total_deaths	new_deaths	total_cases
173	AFG	Asia	Afghanistan	2020-07-01	31517.0	279.0	746.0	13.0	
300	ALB	Europe	Albania	2020-07-01	2535.0	69.0	62.0	4.0	
491	DZA	Africa	Algeria	2020-07-01	13907.0	336.0	912.0	7.0	
613	AND	Europe	Andorra	2020-07-01	855.0	0.0	52.0	0.0	
727	AGO	Africa	Angola	2020-07-01	284.0	8.0	13.0	2.0	

5 rows × 34 columns

- df DataFrame now has one row of data for each country with data present for **July 1st 2020**
- however, it also has a row with a location of World which contains aggregated values for all countries
- df.tail(), df.info() and df.shape will allow for further exploration of the structure of the DataFrame

```
In [2]: ▶ #df.tail()
df.tail()
```

Out[2]:

	iso_code	continent	location	date	total_cases	new_cases	total_deaths	new_deaths	total_ca
29411	ESH	Africa	Western Sahara	2020-07-01	380.0	172.0	1.0	0.0	
29506	YEM	Asia	Yemen	2020-07-01	1158.0	30.0	312.0	8.0	
29623	ZMB	Africa	Zambia	2020-07-01	1594.0	26.0	24.0	2.0	
29738	ZWE	Africa	Zimbabwe	2020-07-01	591.0	17.0	7.0	0.0	
29934	OWID_WRL	NaN	World	2020-07-01	10465987.0	192563.0	511041.0	5732.0	

5 rows × 34 columns

```
In [3]: ▶ #df.info()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 211 entries, 173 to 29934
Data columns (total 34 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   iso_code                               211 non-null    object
 1   continent                              210 non-null    object
 2   location                               211 non-null    object
 3   date                                   211 non-null    datetime64[ns]
 4   total_cases                            210 non-null    float64
 5   new_cases                              210 non-null    float64
 6   total_deaths                           210 non-null    float64
 7   new_deaths                             210 non-null    float64
 8   total_cases_per_million                 210 non-null    float64
 9   new_cases_per_million                   210 non-null    float64
10   total_deaths_per_million                 210 non-null    float64
11   new_deaths_per_million                   210 non-null    float64
12   total_tests                             73 non-null     float64
13   new_tests                               73 non-null     float64
14   total_tests_per_thousand                 73 non-null     float64
15   new_tests_per_thousand                   73 non-null     float64
16   new_tests_smoothed                       83 non-null     float64
17   new_tests_smoothed_per_thousand          83 non-null     float64
18   tests_units                             85 non-null     object
19   stringency_index                         155 non-null    float64
20   population                              211 non-null    float64
21   population_density                       200 non-null    float64
22   median_age                              187 non-null    float64
23   aged_65_old                             184 non-null    float64
24   aged_70_old                             186 non-null    float64
25   gdp_per_capita                           184 non-null    float64
26   extreme_poverty                         122 non-null    float64
27   cvd_death_rate                          186 non-null    float64
28   diabetes_prevalence                     194 non-null    float64
29   female_smokers                           141 non-null    float64
30   male_smokers                             139 non-null    float64
31   handwashing_facilities                   92 non-null     float64
32   hospital_beds_per_thousand               165 non-null    float64
33   life_expectancy                         208 non-null    float64
dtypes: datetime64[ns](1), float64(29), object(4)
memory usage: 57.7+ KB
```

```
In [5]: #df.shape
df.shape
```

```
Out[5]: (211, 34)
```

Q1. Create a new DataFrame called `countries` which is the same as `df` but with the `World` row removed.

- Use the `.copy()` method to ensure you have a distinct DataFrame in memory
- Assign this new DataFrame to the variable `countries` ; do not modify `df`

See below code syntax for some guidance:

```
countries['location'] != 'World'
```

```
In [17]: #add your code below
#countries = df.copy()
#countries = ...

countries = df.copy()
countries = df[df['location'] != 'World'].copy()
countries.head()
```

```
Out[17]:
```

	iso_code	continent	location	date	total_cases	new_cases	total_deaths	new_deaths	total_cases
173	AFG	Asia	Afghanistan	2020-07-01	31517.0	279.0	746.0	13.0	
300	ALB	Europe	Albania	2020-07-01	2535.0	69.0	62.0	4.0	
491	DZA	Africa	Algeria	2020-07-01	13907.0	336.0	912.0	7.0	
613	AND	Europe	Andorra	2020-07-01	855.0	0.0	52.0	0.0	
727	AGO	Africa	Angola	2020-07-01	284.0	8.0	13.0	2.0	

5 rows × 34 columns

Q2. Check the shape of your DataFrame to confirm that `countries` has one row fewer than `df` :

Please note you have been provided with the code for this question to carry out the necessary analysis. Simply uncomment the line of code and run the code cell to produce the desired results.

```
In [12]: #print(df.shape, countries.shape)

print(df.shape, countries.shape)
```

```
(211, 34) (210, 34)
```

Q3. Define a DataFrame based on the `countries` DataFrame, but which only contains the columns in `cols` (defined below) and assign this to a variable called `countries_dr`

- Order this DataFrame by `'total_deaths_per_million'` , with the highest numbers at the top.

See below code syntax for some guidance:

```
DataFrame_name[column_names].sort_values(by=..., ascending=False)
```

```
In [15]: > cols = ['continent', 'location', 'total_deaths_per_million']

#add your code below
#countries_dr

countries_dr = countries[cols].sort_values(by="total_deaths_per_million", ascending=False)
print(countries_dr)
```

	continent	location \
23306	Europe	San Marino
2917	Europe	Belgium
613	Europe	Andorra
28347	Europe	United Kingdom
25362	Europe	Spain
...
23111	North America	Saint Vincent and the Grenadines
23926	Africa	Seychelles
15734	Africa	Lesotho
10808	Europe	Gibraltar
12195	Asia	Hong Kong

	total_deaths_per_million
23306	1237.551
2917	841.615
613	673.008
28347	644.168
25362	606.633
...	...
23111	0.000
23926	0.000
15734	0.000
10808	0.000
12195	NaN

[210 rows x 3 columns]

Q4. Using the countries DataFrame we created earlier, find the sum of total_tests for countries in Africa, assigning the result, as an integer, to africa_tests.

- Use .sum() method calculate the sum for total_tests column
- Use .astype(int) method or int() function to convert results to an integer

See below code syntax for some guidance:

```
countries['continent'] == 'Africa'
```

```
In [27]: > #add your code below
#africa_tests

africa_tests = countries[countries['continent'] == 'Africa']['total_tests'].sum().astype(int)
africa_tests
```

Out[27]: 3445134

Q5. How many countries in Africa have no value recorded for the number of total_tests column? Assign the result to africa_missing_test_data.

- You may find the pandas .isna() method and python len() function useful

See below code syntax for some guidance:

```
len(DataFrame_name[column_name].isna())
```

```
In [116]: africa = countries[countries['continent'] == 'Africa']

africa_tests = africa['total_tests'].sum()

africa_missing_test_data = africa['total_tests'].isna().sum()
africa_missing_test_data
```

Out[116]: 45

Q6. How many countries have a higher value for `total_tests` than the United Kingdom ? Assign your answer to a variable called `countries_more_tests` .

Remember to work from the `countries` DataFrame rather than `df` . You should avoid modifying any existing DataFrames.

```
In [34]: #add your code below
#countries_more_tests

uk_tests = countries[countries['location'] == 'United Kingdom']['total_tests'].values[0]
countries_more_tests = len(countries[countries['total_tests'] > uk_tests])
countries_more_tests
```

Out[34]: 3

Q7. Create a DataFrame called `beds_dr` which is based on the `countries` DataFrame, but contains only the columns `hospital_beds_per_thousand` and `total_deaths_per_million` .

- Your answer should only include rows where there are values present in both of these columns
- You may find the `.dropna()` method useful

See below code syntax for some guidance:

```
DataFrame_name.dropna()
```

```
In [39]: #add your code below
#beds_dr

beds_dr = countries[["hospital_beds_per_thousand", "total_deaths_per_million"]].dropna()
beds_dr
```

Out[39]:

	hospital_beds_per_thousand	total_deaths_per_million
173	0.50	19.163
300	2.89	21.544
491	1.90	20.798
952	3.80	30.635
1081	5.00	28.919
...
29136	0.80	1.794
29332	2.60	0.000
29506	0.70	10.461
29623	2.00	1.305
29738	1.70	0.471

164 rows × 2 columns

Q8. Refer to the `beds_dr` DataFrame. What is the average `total_deaths_per_million` for entries in `beds_dr` where `hospital_beds_per_thousand` is greater than the mean?

- Save the results to a new variable called `dr_high_bed_ratio`

See below code syntax for some guidance:

```
beds_dr['hospital_beds_per_thousand'] > beds_dr['hospital_beds_per_thousand'].mean()
n()
```

```
In [49]: #add your code below
#dr_high_bed_ratio

x = beds_dr['hospital_beds_per_thousand'] > beds_dr['hospital_beds_per_thousand'].mean()
dr_high_bed_ratio = beds_dr[x]['total_deaths_per_million'].mean()
dr_high_bed_ratio
```

Out[49]: 98.18423728813559

Q9. Refer to the `beds_dr` DataFrame. What is the average `total_deaths_per_million` for entries in `beds_dr` where `hospital_beds_per_thousand` is less than the mean?

- Save the results to a new variable called `dr_low_bed_ratio`

See below code syntax for some guidance:

```
beds_dr['hospital_beds_per_thousand'] < beds_dr['hospital_beds_per_thousand'].mean()
n()
```

```
In [50]: #add your code below
#dr_low_bed_ratio

y = beds_dr['hospital_beds_per_thousand'] < beds_dr['hospital_beds_per_thousand'].mean()
dr_low_bed_ratio = beds_dr[y]['total_deaths_per_million'].mean()
dr_low_bed_ratio
```

Out[50]: 56.294057142857135

Q10. Refer to the `countries` DataFrame. Create a new DataFrame called `no_new_cases` which contains only rows from `countries` with zero `new_cases`.

Please note you have been provided with the code for this question to carry out the necessary analysis. Simply uncomment the lines of code and run the code cell to produce the desired results.

```
In [51]: countries.head()
```

Out[51]:

	iso_code	continent	location	date	total_cases	new_cases	total_deaths	new_deaths	total_cases
173	AFG	Asia	Afghanistan	2020-07-01	31517.0	279.0	746.0	13.0	
300	ALB	Europe	Albania	2020-07-01	2535.0	69.0	62.0	4.0	
491	DZA	Africa	Algeria	2020-07-01	13907.0	336.0	912.0	7.0	
613	AND	Europe	Andorra	2020-07-01	855.0	0.0	52.0	0.0	
727	AGO	Africa	Angola	2020-07-01	284.0	8.0	13.0	2.0	

5 rows × 34 columns

```
In [52]: #add your code below
#no_new_cases = countries[countries['new_cases'] == 0]
#no_new_cases.head()

no_new_cases = countries[countries['new_cases'] == 0]
no_new_cases.head()
```

Out[52]:

	iso_code	continent	location	date	total_cases	new_cases	total_deaths	new_deaths	total_cases_
613	AND	Europe	Andorra	2020-07-01	855.0	0.0	52.0	0.0	
836	AIA	North America	Anguilla	2020-07-01	3.0	0.0	0.0	0.0	
952	ATG	North America	Antigua and Barbuda	2020-07-01	66.0	0.0	3.0	0.0	
1381	ABW	North America	Aruba	2020-07-01	103.0	0.0	3.0	0.0	
2080	BHS	North America	Bahamas	2020-07-01	104.0	0.0	11.0	0.0	

5 rows × 34 columns

Q11. Refer to the `no_new_cases` DataFrame. Which country in `no_new_cases` DataFrame has had the highest number of `total_cases` ?

- Save the results to a new variable called `highest_no_new`

See below code syntax for some guidance:

```
no_new_cases['total_cases'] == no_new_cases['total_cases'].max()
```

```
In [117]: highest_no_new = no_new_cases.loc[no_new_cases['total_cases'].idxmax(), 'location']
print(highest_no_new)

Cameroon
```

Q12. Refer to the `countries` DataFrame. What is the sum of the `population` of all countries which have had zero `total_deaths` ?

- Assign your answer to `sum_populations_no_deaths` variable
- Your answer should be in millions, rounded to the nearest whole number, and converted to an integer

```
In [63]: #add your code below
#sum_populations_no_deaths

sum_populations_no_deaths = countries[countries['total_deaths'] == 0]['population'].sum()
sum_populations_no_deaths = round(sum_populations_no_deaths)
sum_populations_no_deaths
```

Out[63]: 192

Q13. Create a function called `country_metric` which accepts the following three parameters:

- a DataFrame (which can be assumed to be of a similar format to `countries`)
- a location (i.e. a string which will be found in the `location` column of the DataFrame)
- a metric (i.e. a string which will be found in any column (other than `location`) in the DataFrame)

The function should return only the value from the first row for a given `location` and `metric` . You may find `.iloc[]` useful.

See below code syntax for some guidance:

```
def country_metric(df, location, metric):  
  
    return df[df['location'] == location].iloc[0][metric]
```

```
In [65]: ▶ #add your code below  
#def country_metric(df, location, metric):  
  
def country_metric(df, location, metric):  
    return countries[countries["location"] == location].iloc[0][metric]
```

Q.14 Use your function to collect the value for Vietnam for the metric aged_70_older , assigning the result to vietnam_older_70 .

Please note you have been provided with the code for this question to carry out the necessary analysis. Simply uncomment the lines of code and run the code cell to produce the desired results.

```
In [67]: ▶ #add your code below  
#vietnam_older_70 = country_metric(countries, 'Vietnam', 'aged_70_older')  
#vietnam_older_70  
  
vietnam_older_70 = country_metric(countries, 'Vietnam', 'aged_70_older')  
vietnam_older_70
```

Out[67]: 4.718

Q.15 Create another function called countries_average , which accepts the following three parameters:

- a DataFrame "df" (which can be assumed to be such as countries)
- a list of countries "countries" (which can be assumed to all be found in the location column of the DataFrame)
- a string "metric" (which can be assumed to be a column (other than location) which will be found in the DataFrame) . For instance, this string value can be life_expectancy .

Note that for the test on KATE for this question to pass, you need to make sure the function accepts the three parameters in the following order: countries_average(df, countries, metric) . (You can call your parameters however you like as long as the type of these parameters are what was described above).

The function should return the average value for the given metric for the given list of countries.

You may find .isin() method useful while filtering for list of countries.

```
In [106]: ▶ #add your code below  
#def countries_average(df, countries, metric):  
  
def countries_average(df, countries, metric):  
    new_df = df[df['location'].isin(countries)]  
  
    average_value = new_df[metric].mean()  
  
    return average_value
```

Q16. Use your countries_average function to find out the average life_expectancy of countries in the g7 list defined below. Assign the result to the variable g7_avg_life_expectancy .

Please note you have been provided with the code for this question to carry out the necessary analysis.
Simply uncomment the lines of code and run the code cell to produce the desired results.

```
In [107]: g7 = ['United States', 'Italy', 'Canada', 'Japan', 'United Kingdom', 'Germany', 'France']
```

```
In [108]: countries.head()
```

Out[108]:

	iso_code	continent	location	date	total_cases	new_cases	total_deaths	new_deaths	total_cases
173	AFG	Asia	Afghanistan	2020-07-01	31517.0	279.0	746.0	13.0	
300	ALB	Europe	Albania	2020-07-01	2535.0	69.0	62.0	4.0	
491	DZA	Africa	Algeria	2020-07-01	13907.0	336.0	912.0	7.0	
613	AND	Europe	Andorra	2020-07-01	855.0	0.0	52.0	0.0	
727	AGO	Africa	Angola	2020-07-01	284.0	8.0	13.0	2.0	

5 rows × 34 columns

```
In [109]: #add your code below
g7 = ['United States', 'Italy', 'Canada', 'Japan', 'United Kingdom', 'Germany', 'France']
g7_avg_life_expectancy = countries_average(df, g7, 'life_expectancy')
g7_avg_life_expectancy
```

Out[109]: 82.10571428571428

Q.17 Refer to the `countries` DataFrame. Find the country with lowest value for `life_expectancy` in the `countries` DataFrame, and create a string which is formatted as follows:

'{country} has a life expectancy of {diff} years lower than the G7 average.'

Assign your string to the variable `headline` and ensure it is formatted exactly as above, with:

- use f-strings to format the string
- {country} being replaced by the value in the `location` column of the DataFrame
- {diff} being replaced by a float **rounded to one decimal place**, of the value from the `life_expectancy` column subtracted from `g7_avg_life_expectancy`. Please note that {diff} should be a positive value

```
diff = <G7 countries average life expectancy> - <value of the lowest life expectancy country>
```

See below code syntax for some guidance:

```
lowest = countries[countries['life_expectancy'] == countries['life_expectancy'].min()].iloc[0]
country = lowest['location']
life_exp = lowest['life_expectancy']
```

```
In [110]: ► #add your code below
#headline = f'{country} has a life expectancy of {diff} years lower than the G7 average'

lowest = countries[countries['life_expectancy'] == countries['life_expectancy'].min()]
country = lowest['location']
life_exp = lowest['life_expectancy']

diff = g7_avg_life_expectancy - life_exp

headline = f"{country} has a life expectancy of {diff:.1f} years lower than the G7 average"
print(headline)

Central African Republic has a life expectancy of 28.8 years lower than the G7 average.
```

In []: ►